

Planning in the Taxi Domain

Overview

This exercise involves modeling a sequential decision-making problem. We consider the Taxi Domain problem that involves controlling a taxi agent that can pick up and drop passengers in a grid world. We will implement algorithms for arriving at a good policy offline using dynamic programming and later implement basic reinforcement learning methods for online learning. You are requested to submit your implementation and a report that includes the insights gained via this exercise.

The Taxi Domain

- a. A Taxi agent is situated in a 5x5 grid world (see figure below). The goal of the taxi agent is to pick up a passenger from an initial grid cell and drop the passenger at a destination grid cell. There are four specially-designated locations in this grid world, called *depots*, that are denoted as R, G, B and Y (initials for Red, Green, Blue and Yellow). The taxi agent is shown using a curved symbol. There are a few walls situated along grid boundaries (indicated by thick black lines) which prevent the taxi from moving across grid cells. The entire grid is assumed to be enclosed by walls along the left, right, top and bottom sides.




Figure: The Taxi Domain (5x5 grid world). The taxi is indicated by a curved symbol. There are 4 depots indicated as R, G, B and Y. The taxi agent aims at transporting passengers from the source depot to the destination depot.

- b. There are six actions in the domain: (a) four *navigation* actions that move the taxi one grid cell to the *North*, *South*, *East* or the *West* directions, (b) a *Pickup* action, where the taxi attempts to pick up the passenger and (c) a *Putdown* action, where the taxi drops off the passenger. Each navigation action succeeds with a probability of 0.85 and moves in a

random other direction with probability 0.15. Movements that attempt to cross the boundary of the grid or a wall results in no change in the state. The *Pickup* and the *Putdown* actions are deterministic and lead to their exact intended effects.

- c. The taxi agent receives a reward of (-1) for performing each action. A reward of (+20) is received when the taxi agent successfully delivers the passenger at the destination grid cell. Further, a reward of (-10) is received if the taxi attempts to *Pickup* or *Putdown* a passenger when the taxi and the passenger are not located in the same grid cell.
- d. The taxi agent interacts in the environment in episodes. In each episode, the taxi starts in a randomly chosen grid cell. There is a passenger at one of the four depot locations chosen randomly and that passenger wishes to be transported to one of the other depot locations. The destination is different from the source and also selected randomly. The taxi must move towards the passenger's grid cell (called the source), pick up the passenger, go to the destination location (called the destination), and drop the passenger there. The episode ends when the passenger is deposited at the destination location. Note that dropping the passenger in a location *other* than the destination will not terminate the episode.

Part A: Computing Policies

In this section, we model the taxi domain as a Markov Decision Process (MDP) and compute optimal policy for the taxi agent.

1. Formulate the taxi domain as an MDP
 - a. Describe the state space, the action space, the transition model and the reward model for the problem in the report.
 - b. Implement a simulator for the taxi domain which includes a model of the environment and determines the next state based on the action taken by the agent. The stochastic effects of navigation actions are simulated as part of the environment model. The agent should receive an instantaneous reward after taking each action.
 - c. An instance of the taxi domain problem consists of a starting depot for the passenger, selecting a different destination depot and selecting a starting location for the taxi in the grid.
2. Implement Value Iteration for the taxi domain.
 - a. The implementation should take as input an instance of the taxi domain and a parameter ϵ (denoting the maximum error allowed in the value of any state) and return the optimal policy. Use the *max-norm* distance¹ in the successive value functions to determine convergence. Initially, set the discount factor as 0.9. Report the value of ϵ you have chosen and the number of iterations required for the convergence in your write-up.
 - b. Next, we study the connection between the discount factor and the rate of convergence. Repeat part (a) by varying the discount factor in the range $\{0.01, 0.1, 0.5, 0.8, 0.99\}$. For each discount factor considered, plot the iteration index on the x-axis and the *max-norm* distance along the y-axis. Describe your observations.
 - c. Pick an instance such as follows: Y (passenger initial location), G (passenger destination location) and R (taxi initial location). Simulate the policy obtained for the discount factor γ as 0.1 and 0.99. Determine the first 20 states and actions prescribed by the policy. Examine the two state-action sequences. Report any differences in the execution traces for the policies obtained for the two discount factors. Repeat the runs by keeping the goal the same and varying the start states (for the taxi and the passenger).

¹ Reference: AIMA Chapter 17.

3. Implement Policy Iteration for the problem
 - a. How would you implement the Policy Evaluation step? Implement a linear algebra method and an iterative method for Policy Evaluation and report when one approach would be better than the other.
 - b. Run policy iteration till convergence to determine the optimal policy. Repeat the process by computing the policy loss at each iteration between the policy at each iteration the final policy at convergence. Plot the policy loss (along y-axis) and iteration number (x-axis). Study the behaviour by varying the discount factor in the range $\{0.01, 0.1, 0.5, 0.8, 0.99\}$.

Part B: Incorporating Learning

In this section, assume that we have an *unknown* taxi domain MDP where the transition model and the reward model is not known to the taxi agent. In this case, the taxi must act only based on its past experience. Assume that the next state and reward is provided by the simulated environment when the taxi agent executes an action.

1. Implement the following (model-free) approaches to learn the optimal policy.
Implement
 - a. Implement Q-learning that uses a 1-step greedy look-ahead policy for the Q-learning updates. Incorporate an epsilon-greedy policy for exploration with a fixed exploration rate (epsilon) of 0.1.
 - b. Implement Q-learning with an epsilon-greedy policy with a decaying exploration rate. The exploration is initialized with epsilon as 0.1 and is decayed, inversely proportional to the iteration number. Here, a single iteration refers to a single Q-learning update.
 - c. Implement *State-Action-State-Action-Reward* (SARSA) for this problem. Use an epsilon-greedy policy for exploration with a fixed exploration rate (epsilon) of 0.1.
 - d. As in part (b) implement SARSA with a decaying exploration rate.
2. Execute the above algorithms for *at least* 2000 episodes. Initialize the above methods with a learning rate (alpha) as 0.25 and use a discount factor (gamma) of 0.99. Assume that each episode begins from a randomly selected feasible state and terminates if the taxi agent reaches the goal state or the episode reaches a maximum length of 500 steps. Evaluate each algorithm by computing the sum of discounted rewards accumulated during the episode (averaged over 10 different runs). Plot the sum of discounted rewards accumulated during an episode (along y-axis) against the number of training episodes (along x-axis). Analyse the convergence of the algorithms.
3. From part 2, select the learning algorithm that converges to the highest accumulated reward. Execute the policy on 5 problem instances by varying the initial depot location of the passenger and the initial depot location of the taxi agent. What do you observe?
4. Consider the Q-learning method implemented in 1a. Plot the sum of discounted rewards accumulated during an episode against the number of training episodes by varying the exploration rate (epsilon) as {0, 0.05, 0.1, 0.5, 0.9} with a learning rate (alpha) kept as 0.1. Analyze the impact of using a high or a low exploration rate. Next, vary the learning rate (alpha) as {0.1, 0.2, 0.3, 0.4, 0.5} keeping the exploration rate (epsilon) as 0.1.
5. Consider the 10x10 extension of the taxi domain problem (see figure below). Select the best learning algorithm from your experiments above and try the learner on the 10x10 problem. Assume a maximum of 10,000 learning iterations. Report the accumulated discounted reward for the learned policy from your method. Average over at least 5

instances sampling the initial states for the passenger and the taxi and the destination location for the taxi among the depots.




Figure: The 10x10 version of the Taxi Domain. The taxi is indicated by a curved symbol. There are 8 depots indicated as R, G, B, Y, C, P, W, M. The taxi agent aims at transporting passengers from source to destination locations.